



# Deploying JupyterHub for students and researchers

JupyterCon 2017

Min Ragan-Kelley, Simula  
Carol Willing, Cal Poly  
Yuvi Panda, UC Berkeley  
Ryan Lovett, UC Berkeley

# Set Up

```
git clone \  
    https://github.com/jupyterhub/jupyterhub-tutorial \  
    /srv/jupyterhub
```



# Tutorial logistics

9:00 - 9:15      Welcome

9:15 - 10:30    JupyterHub

10:30 - 11:00   Morning break

11:00 - 12:20   JupyterHub and Kubernetes

12:20 - 12:30   Wrap up



# Tutorial logistics

- Gitter channel for this tutorial

<https://gitter.im/jupyterhub-tutorial/Lobby>

- Post-its

questions and assistance

done with task and doing well

- Tutorial materials

<https://github.com/jupyterhub/jupyterhub-tutorial>





# What are Jupyter and JupyterHub?





# What is a Notebook?

- Document
- Environment
- **Web app**

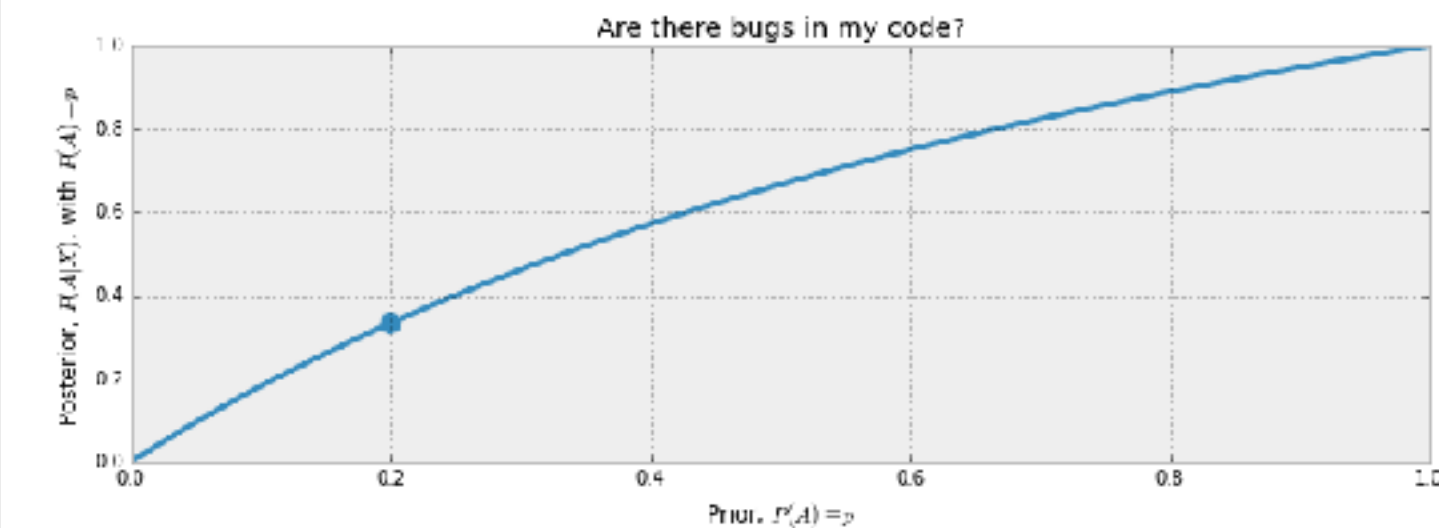
We have already computed  $P(X|A)$  above. On the other hand,  $P(X| \sim A)$  is subjective: our code can pass tests but still have a bug in it, though the probability there is a bug present is reduced. Note this is dependent on the number of tests performed, the degree of complication in the tests, etc. Let's be conservative and assign  $P(X| \sim A) = 0.5$ . Then

$$P(A|X) = \frac{1 \cdot p}{1 \cdot p + 0.5(1 - p)}$$
$$= \frac{2p}{1 + p}$$

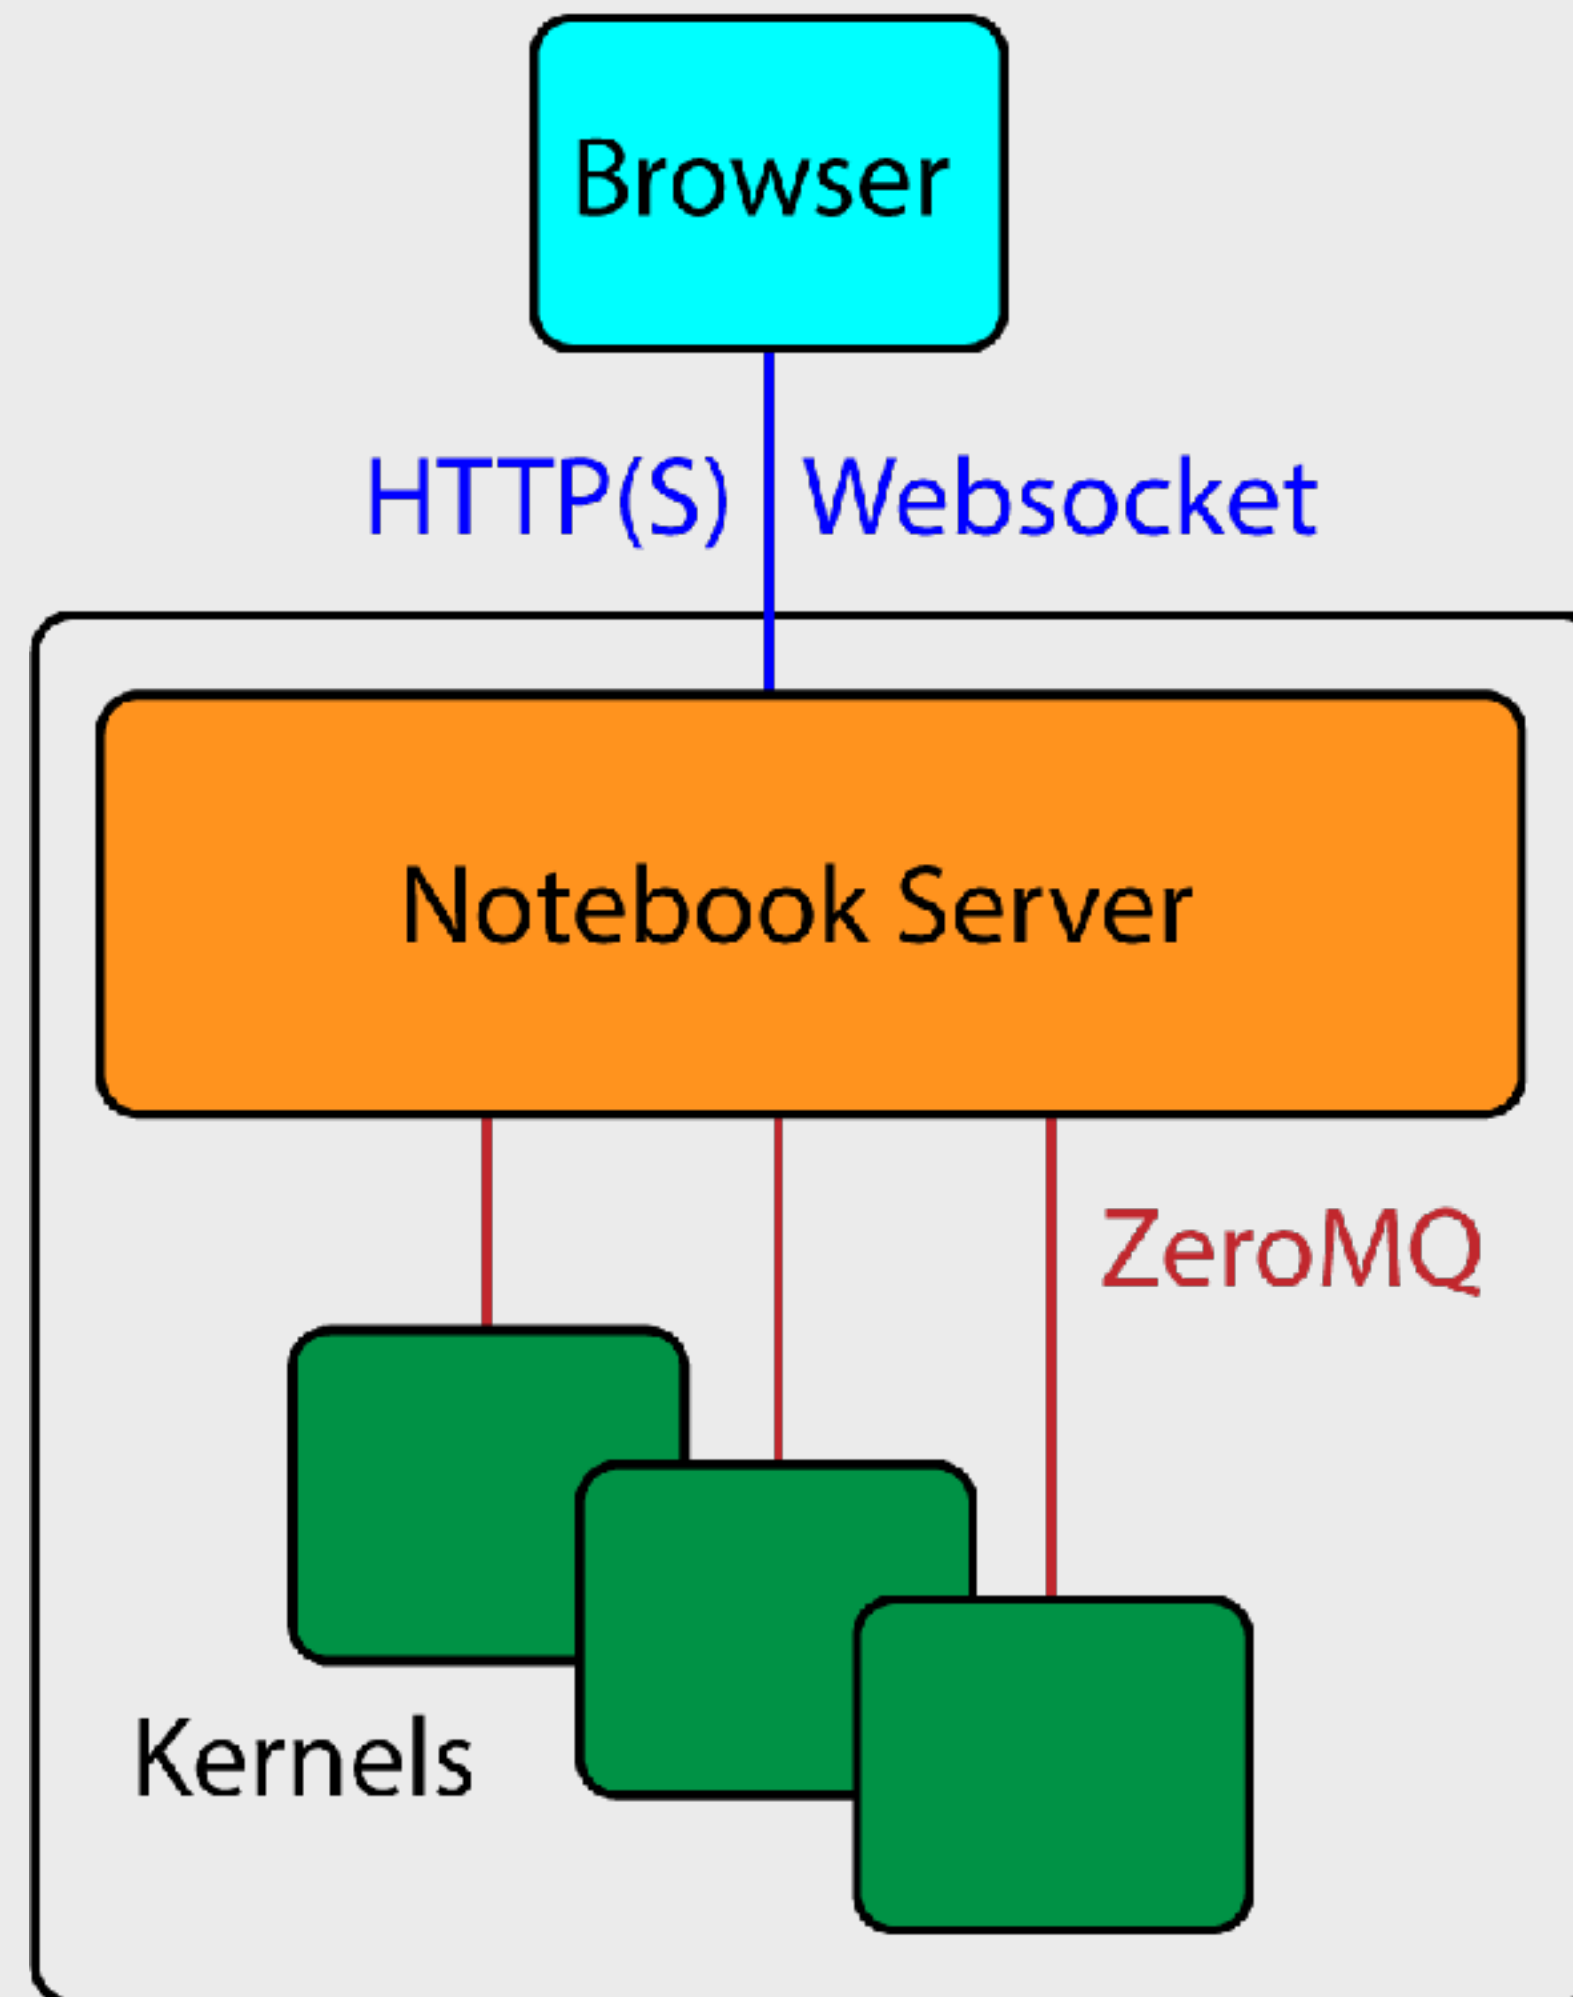
This is the posterior probability. What does it look like as a function of our prior,  $p \in [0, 1]$ ?

```
figsize(12.5, 4)
p = np.linspace(0, 1, 50)
plt.plot(p, 2 * p / (1 + p), color="#348ABD", lw=3)
# plt.fill_between(p, 2*p/(1+p), alpha=.5, facecolor="#A68628")
plt.scatter(0.2, 2 * (0.2) / 1.2, s=140, c="#348ABD")
plt.xlim(0, 1)
plt.ylim(0, 1)
plt.xlabel("Prior, $P(A) = p$")
plt.ylabel("Posterior, $P(A|X)$, with $P(A) = p$")
plt.title("Are there bugs in my code?")
```

<matplotlib.text.Text at 0x1051de550>

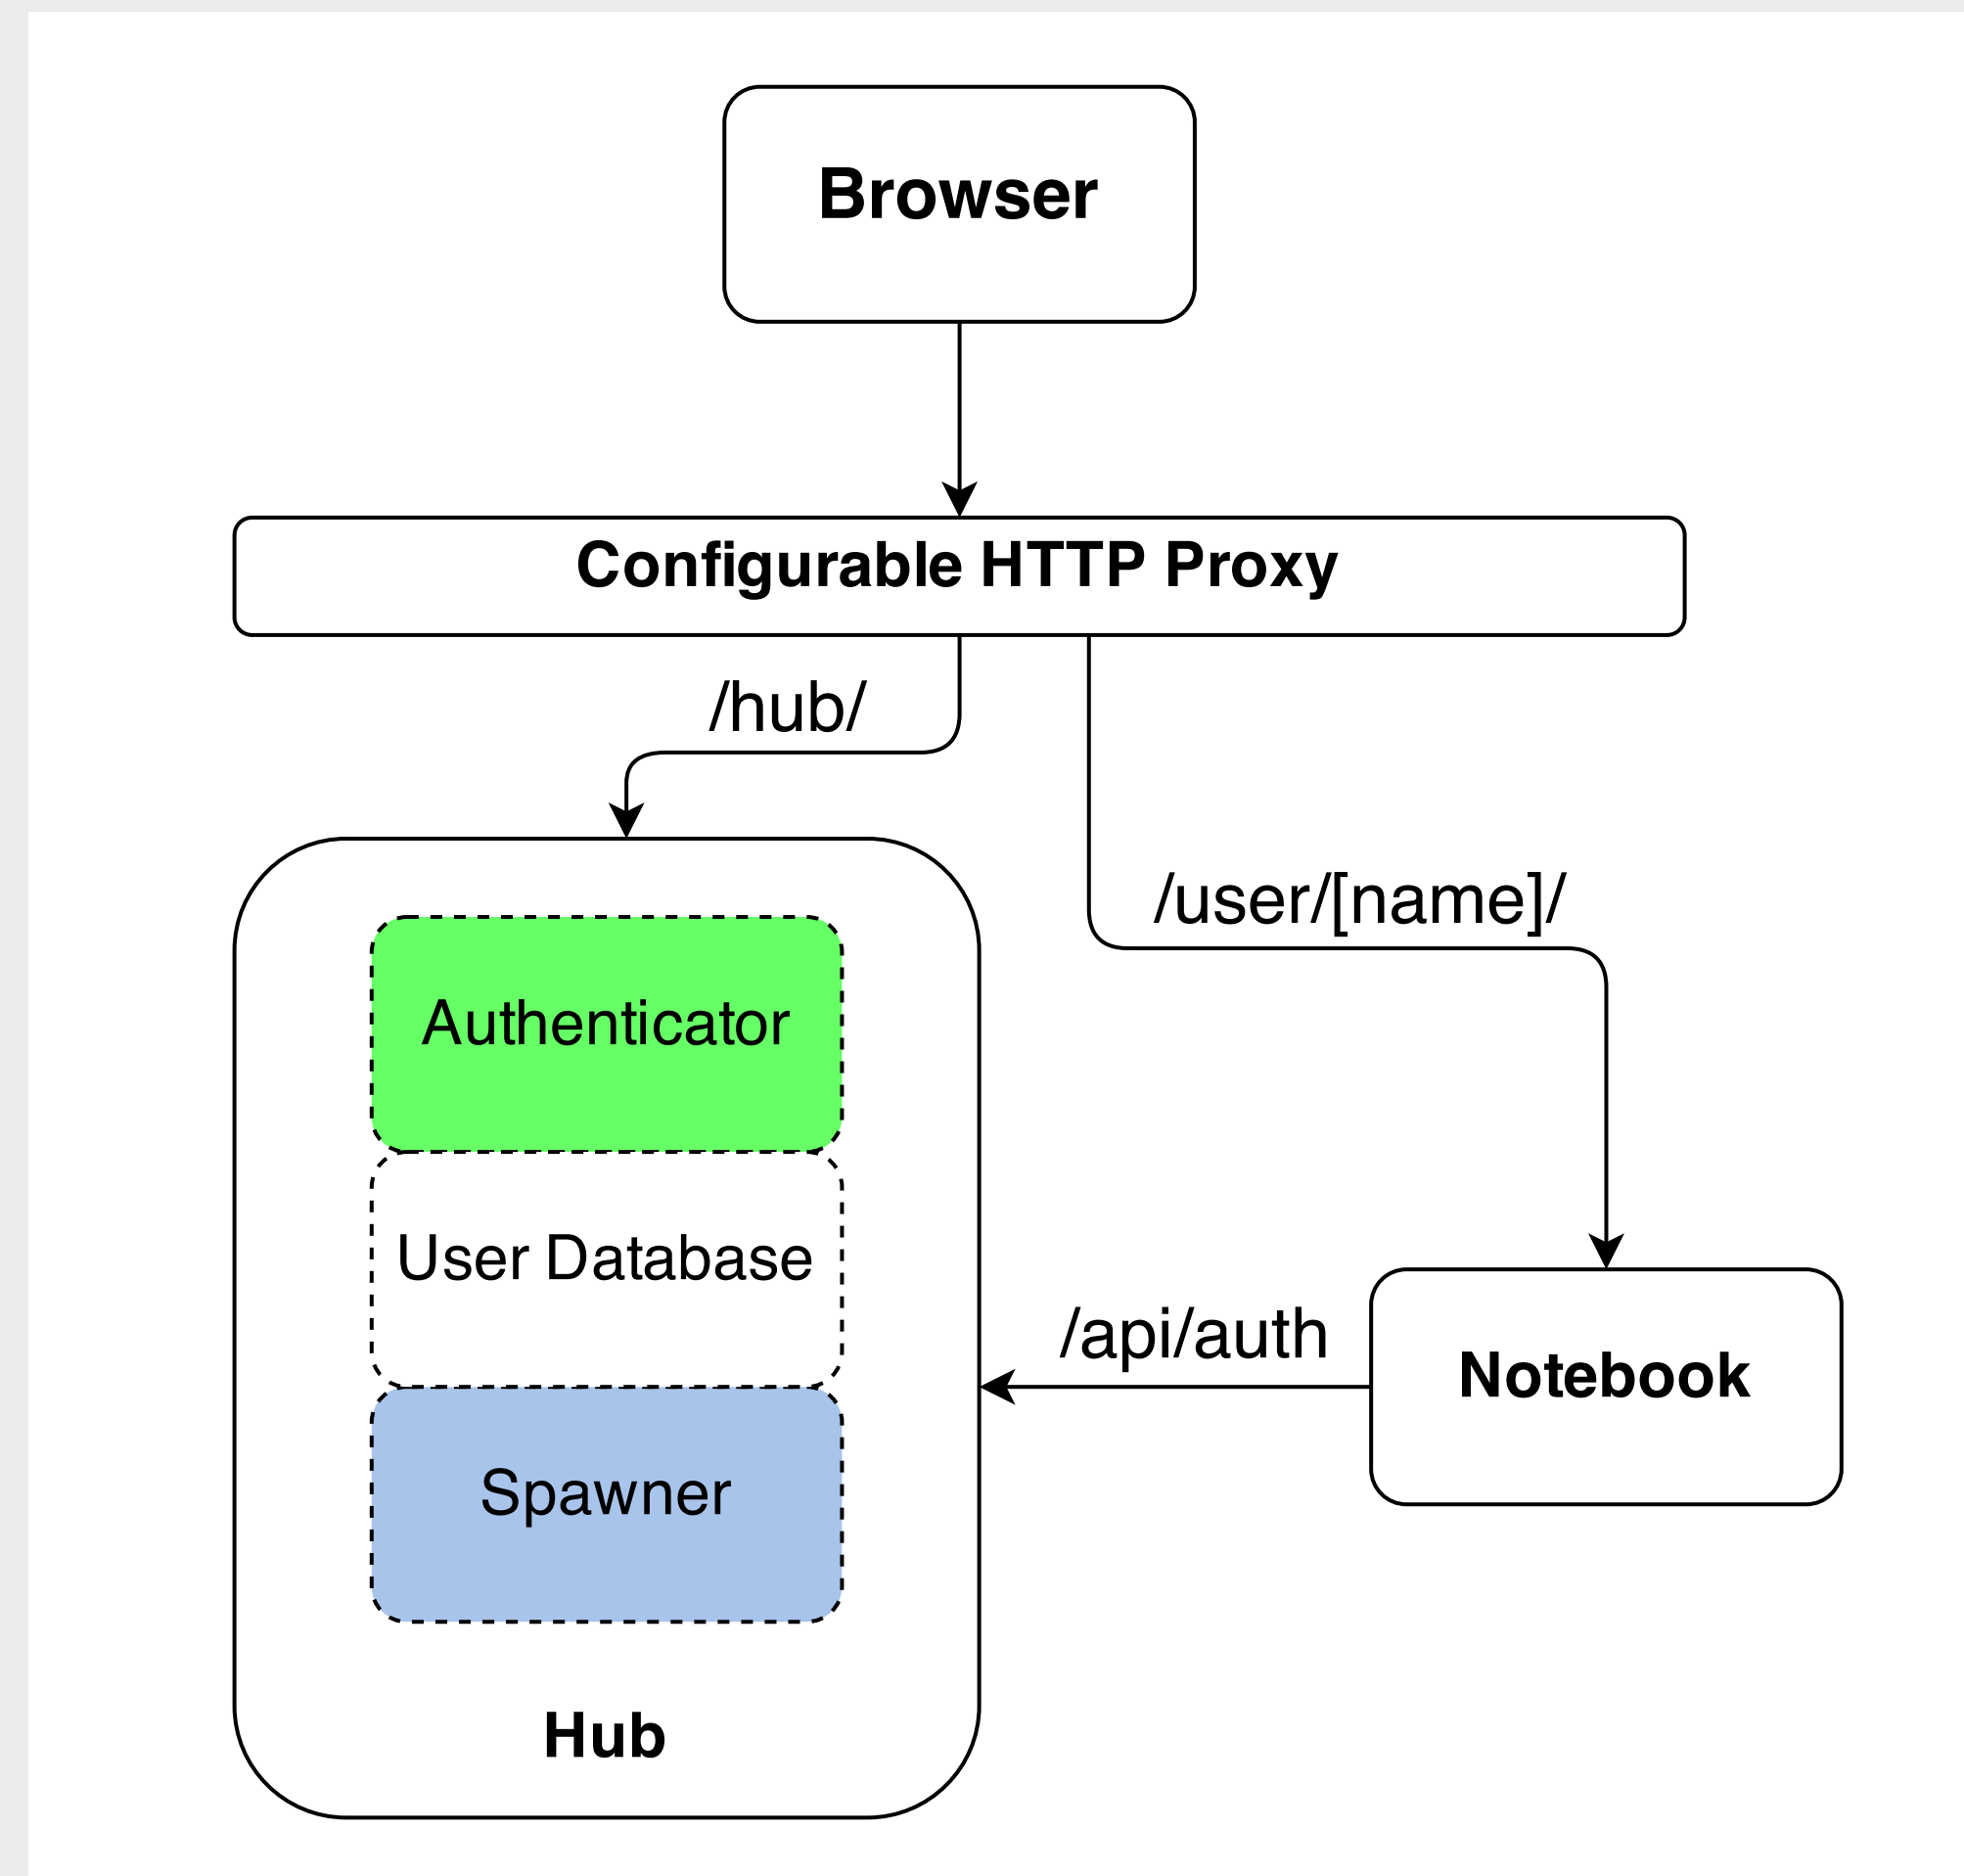


# What is a Notebook Server?



# jupyterhub

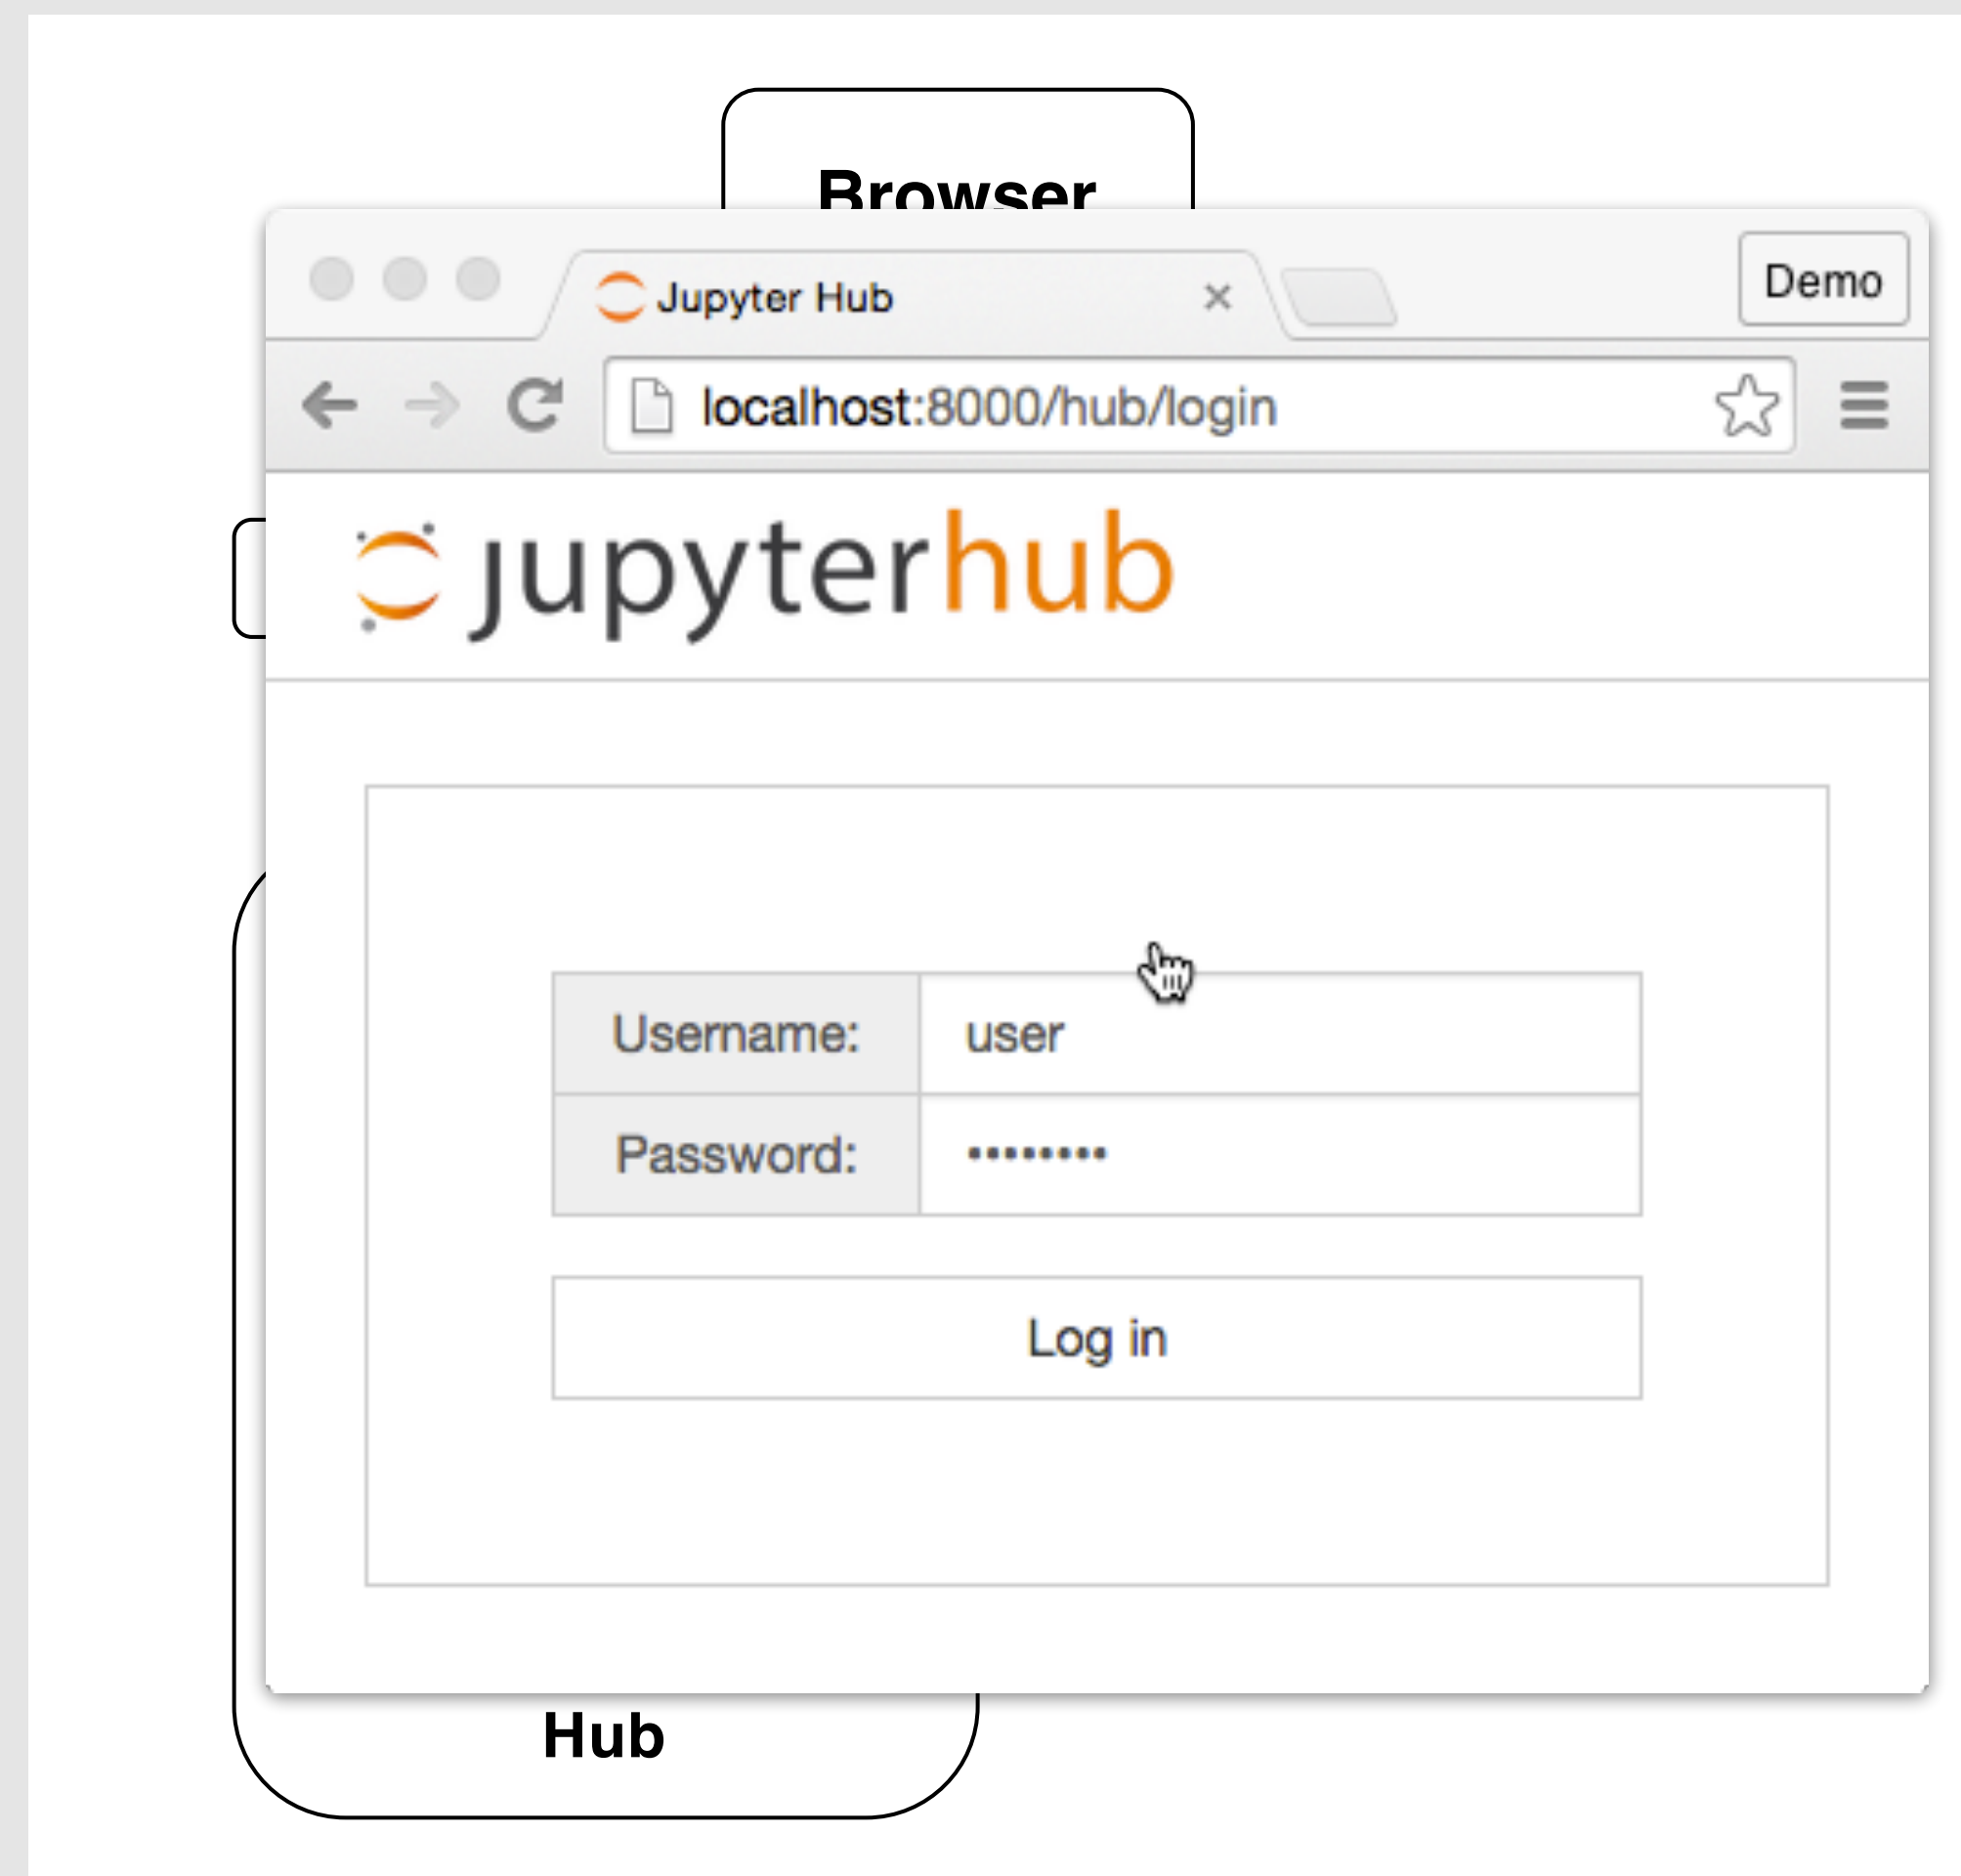
- Manages authentication
- Spawns single-user servers on-demand
- Each user gets a complete notebook server





# jupyterhub

- Initial request is handled by Hub
- User authenticates via form / OAuth
- Spawner starts single-user server
- Hub notifies Proxy
- Redirects user to `/user/[name]`
- Single-user Server verifies auth with Hub





# Installation



*Photo taken by Matthew Bowers*





# Installation (as administrator)

conda:

```
conda install -c conda-forge jupyterhub  
conda install notebook
```

pip, npm:

```
python3 -m pip install jupyterhub  
npm install -g configurable-http-proxy
```

test:

```
jupyterhub -h  
configurable-http-proxy -h
```



# Installation (this repo)

```
conda env create -f environment.yml  
source activate jupyterhub-tutorial
```





# Installation Caveats

JupyterHub installation must be  
**readable+executable** by all users\*

This is often not the case for envs, so be  
careful

\*when using local users



# Plug: conda-forge



Community-managed conda packages.

<https://conda-forge.github.io>

```
conda config --add channels conda-forge
```

# Installation: Spawner and Single User

<https://docs.docker.com/engine/installation>

```
pip install dockerspawner  
docker pull jupyterhub/singleuser
```



# JupyterHub Defaults

- Authentication: PAM (local users, passwords)
- Spawning: Local users
- Hub must run as root





# Aside: SSL

- JupyterHub is an authenticated service - users login. That should **never** happen over plain HTTP.
- For testing, we can generate self-signed certificates:

```
openssl req -x509 -nodes -days 365 -newkey rsa:1024 \
    -keyout jupyterhub.key -out jupyterhub.crt
```

Note: Safari will not connect websockets to untrusted (self-signed) certs



# SSL: Let's Encrypt

- <https://letsencrypt.org/getting-started/>
- Free SSL for any domain

```
wget https://dl.eff.org/certbot-auto
```

```
chmod a+x certbot-auto
```

```
./certbot-auto certonly --standalone -d mydomain.tld
```

```
key: /etc/letsencrypt/live/mydomain.tld/privkey.pem
```

```
cert: /etc/letsencrypt/live/mydomain.tld/fullchain.pem
```





# Configuration



*Photo taken by Matthew Bowers*





# Start configuring JupyterHub

```
jupyterhub --generate-config
```

```
c.JupyterHub.ssl_key = 'jupyterhub.key'  
c.JupyterHub.ssl_cert = 'jupyterhub.crt'  
c.JupyterHub.port = 443
```

```
jupyterhub_config.py
```





# Installing language kernels for all users

```
conda create -n py2 python=2 ipykernel  
conda run -n py2 -- ipython kernel install  
  
jupyter kernelspec list
```



# Using GitHub OAuth

<https://github.com/settings/applications/new>

**Register a new OAuth application**

**Application name**  
  
Something users will recognize and trust


**Homepage URL**  
  
The full URL to your application homepage

**Application description**  
  
This is displayed to all potential users of your application

**Authorization callback URL**  
  
Your application's callback URL. Read our [OAuth documentation](#) for more information.

**Register application**

**JupyterHub tutorial**

 **minrk** owns this application. [Transfer ownership.](#)

**1 user**

**Client ID**  
862e21790aeb654a0c98

**Client Secret**  
c2c1 [REDACTED] befa6

**Revoke all user tokens**

**Reset client secret**



# Using GitHub OAuth

File: `./env`

```
export GITHUB_CLIENT_ID=from_github  
export GITHUB_CLIENT_SECRET=from_github  
export OAUTH_CALLBACK_URL=https://YOURDOMAIN/hub/oauth\_callback
```

```
source ./env
```



# Using GitHub OAuth

We need to install: OAuthenticator

```
python3 -m pip install oauthenticator
```

Config file: jupyterhub\_config.py

```
from oauthenticator.github import LocalGitHubOAuthenticator
c.JupyterHub.authenticator_class = LocalGitHubOAuthenticator
c.LocalGitHubOAuthenticator.create_system_users = True
```





# Specifying Users

By default, any user that successfully authenticates is allowed to use the Hub.

This is appropriate for shared workstations with PAM Auth, but probably not GitHub:

```
# set of users allowed to use the Hub  
c.Authenticator.whitelist = {'minrk', 'takluyver'}  
  
# set of users who can administer the Hub itself  
c.Authenticator.admin_users = {'minrk'}
```





# Custom Authenticators



*Photo taken by Matthew Bowers*





# Using DockerSpawner

We need DockerSpawner:

```
python3 -m pip install dockerspawner netifaces  
docker pull jupyterhub/singleuser
```

In `jupyterhub_config.py`:

```
from oauthenticator.github import GitHubOAuthenticator  
c.JupyterHub.authenticator_class = GitHubOAuthenticator  
  
from dockerspawner import DockerSpawner  
c.JupyterHub.spawner_class = DockerSpawner
```



# Using DockerSpawner

```
from dockerspawner import DockerSpawner
c.JupyterHub.spawner_class = DockerSpawner

# The Hub's API listens on localhost by default,
# but docker containers can't see that.
# Tell the Hub to listen on its docker network:
import netifaces
docker0 = netifaces.ifaddresses('docker0')
docker0_ipv4 = docker0[netifaces.AF_INET][0]
c.JupyterHub.hub_ip = docker0_ipv4['addr']
```



# Using DockerSpawner

- There is *\*loads\** to configure with Docker
- Networking configuration
- Data volumes
- `DockerSpawner.image = 'jupyter/scipy-notebook:8f56e3c47fec'`





# Customizing Spawners



*Photo taken by Matthew Bowers*





# JupyterHub with supervisor

apt-get install supervisor

```
#!/usr/bin/env bash
# /srv/jupyterhub/launch.sh
set -e
source ./env
exec jupyterhub $@
```

```
# /etc/supervisor/conf.d/jupyterhub.conf
[program:jupyterhub]
command=bash launch.sh
directory=/srv/jupyterhub
autostart=true
autorestart=true
startretries=3
exitcodes=0,2
stopsignal=TERM
redirect_stderr=true
stdout_logfile=/var/log/jupyterhub.log
stdout_logfile_maxbytes=1MB
stdout_logfile_backups=10
user=root
```



# Reference Deployments

<https://github.com/jupyterhub/helm-chart>

Helm, KubeSpawner, Hub in container

<https://github.com/jupyterhub/jupyterhub-deploy-docker>

docker-compose, DockerSpawner, Hub in Docker

<https://github.com/jupyterhub/jupyterhub-deploy-teaching>

ansible, no docker, nbgrader



# Docker Deployment

- Docker Compose: <https://docs.docker.com/compose/install/>
- git clone <https://github.com/jupyterhub/jupyterhub-deploy-docker>
- Setup the basics (creates volumes, network, pulls images):  
    make build



# Docker Deployment

- `mkdir secrets`
- Copy SSL key, cert to:
  - `secrets/jupyterhub.cert` (cert)
  - `secrets/jupyterhub.key` (key)



# Docker Deployment

Make userlist:

```
minrk admin  
willingc admin  
yuvipanda  
ryanlovett
```





# Docker Deployment

Launch: 

```
docker-compose up
```





# Optimizations and Best Practices



*Photo taken by Matthew Bowers*





# Optimizations and best practices

- **Always use SSL!**
- Use postgres for the Hub database
- Put nginx in front of the proxy
- Run cull-idle-servers service to prune resources
- Global configuration in `/etc/jupyter` and `/etc/ipython`
- **Back up your user data!!!**



# When to use JupyterHub

- A class where students can do homework (nbgrader)
- A short-lived workshop, especially if installation is hard
- A research group with a shared workstation or small cluster
- On-site computing resources for researchers and analysts at an institution



# When *not* to use JupyterHub

- JupyterHub is Authenticated and Persistent
- It takes work to keep it going
- SageMathCloud is *hosted* and provides realtime-collaboration
- A growing number of cloud providers support notebook services (Google, Microsoft, etc.)





# JupyterHub API



*Photo taken by Matthew Bowers*





# End of JupyterHub Fundamentals

Break - 30 min

See you back here at 11:00 am.

Up next:  
JupyterHub and  
Kubernetes



# Attributions and recognition

A huge thank you to the Project Jupyter team and community.  
Your hard work and passion makes this all possible.

GORDON AND BETTY  
**MOORE**  
FOUNDATION



Alfred P. Sloan  
FOUNDATION

THE LEONA M. AND HARRY B.  
**HELMSLEY**  
CHARITABLE TRUST





# Thank you





# Thank you

*try.jupyter.org*

*www.jupyter.org*

*ipython.org*

*numfocus.org*

